



Filesystems and Storage on RMACC Summit

Filesystems and Storage on RMACC Summit

- Daniel Trahan
 - *Email:* Daniel.Trahan@Colorado.edu
 - *RC Homepage:* <https://www.colorado.edu/rc>
 - *RC Email:* rc-help@colorado.edu
-
- Slides available for download at:
 - https://github.com/ResearchComputing/Filesystems_And_Storage_Fall_2020

Outline

- Overview of Summit's storage
- Summit architecture and filesystems
- Petalibrary
- Data transfers and tools

Quick note

- Clusters come in all shapes and sizes, so much of the information here may not apply to other systems. HPC is a very diverse landscape so make sure you check with the system administrators of whichever cluster you are using.

Overview of RC's directories

- 3 major user directories
 - Home – Used for reusable job scripts, setting files, and other important small files.
 - Projects – Used for application and small datasets.
 - Scratch – Work directory. Used with jobs for highspeed access to data or output.
- Table:

	Directory	Capacity	Backup	Purge
Home	/home/\$USER	2 GB	2 hours for 7 days	Never
Projects	/projects/\$USER	250 GB	6 hours for 7 days	Never
Scratch	/scratch/summit/\$USER	10 TB	(none)	90 days

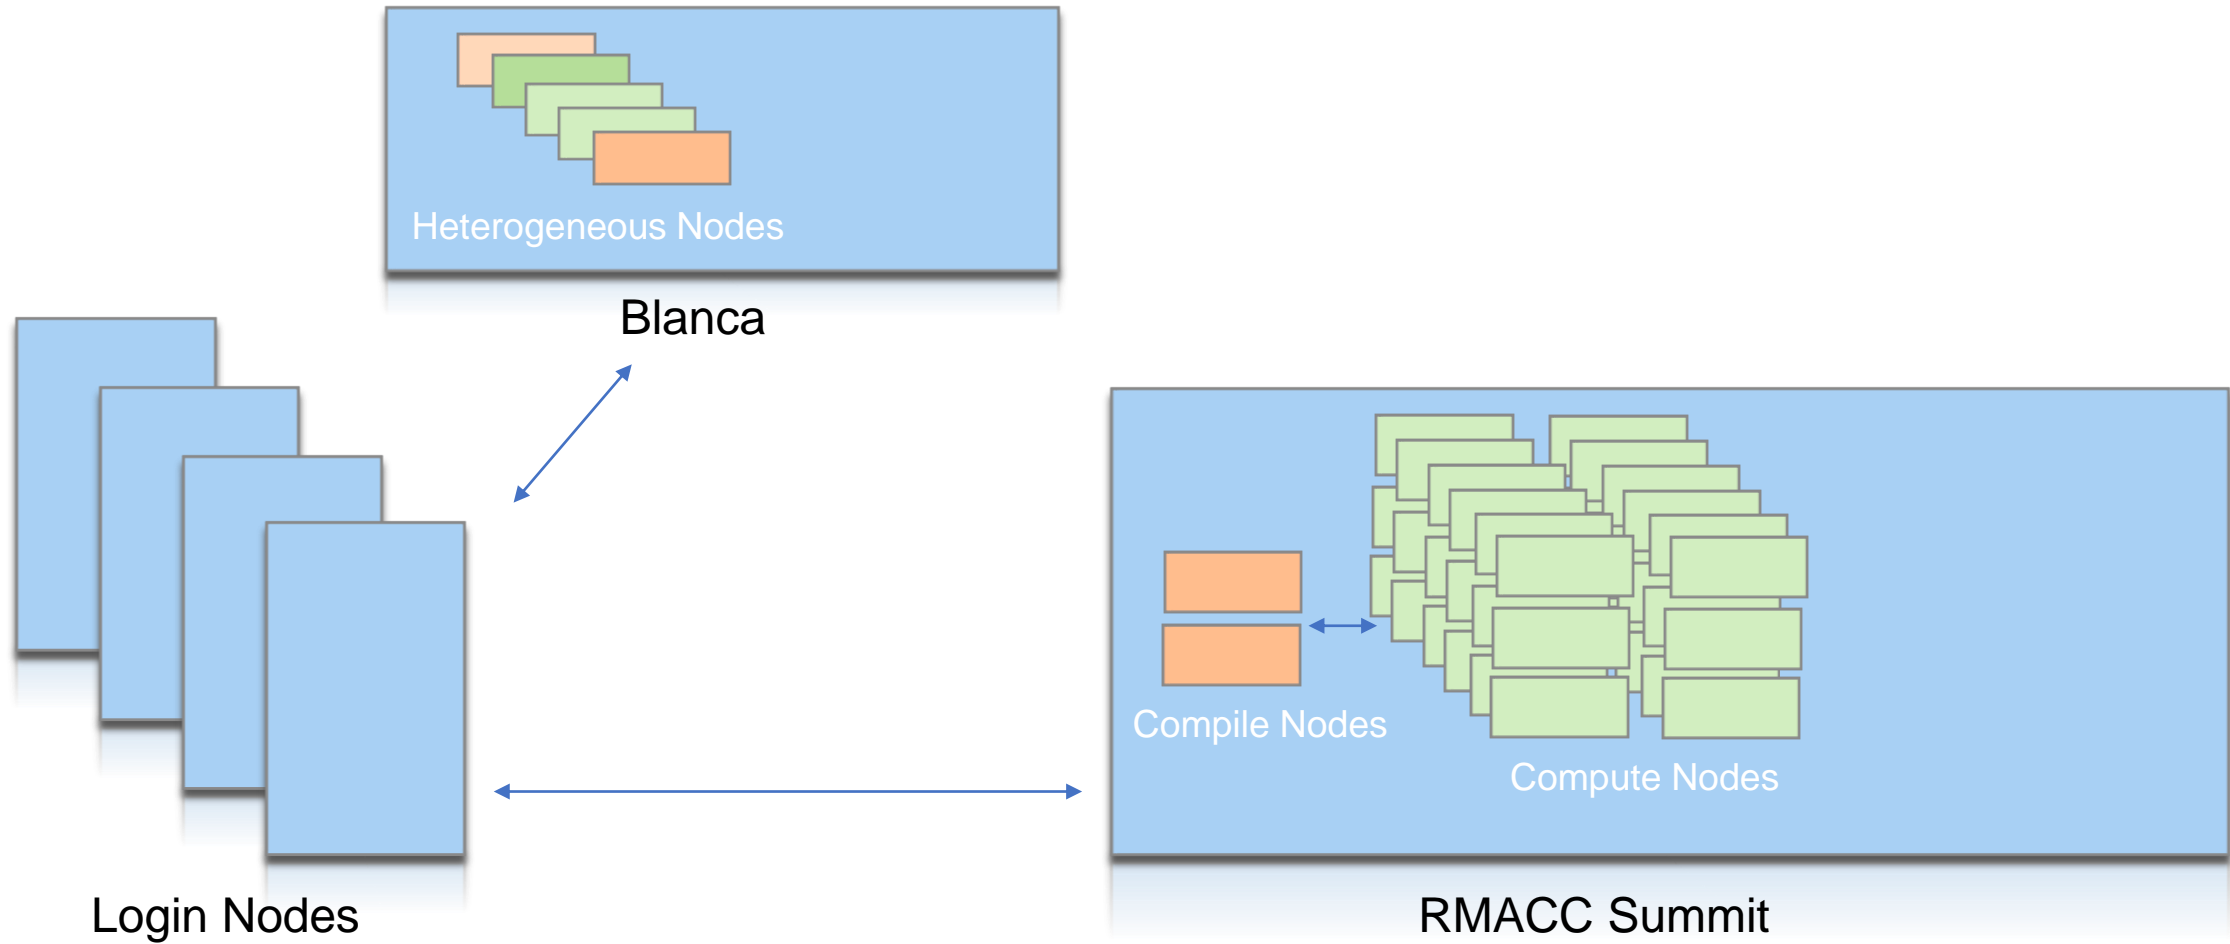
What's makes Summit different?

- As you may know Research Computing resources diverge a bit from the average computer.
 - Complex servers architecture connected to two separate clusters
 - Various endpoints and services
 - Several Filesystems for various purposes.
- When you log into Summit all of this seems “seamless...”
- So what gives?

RC's network

- The user facing side of RC's network is composed of 3 key compute structures:
- Login Nodes
 - VMs with very little compute capability
 - Can access RMACC Summit by running `ssh scompile` with the `slurm/summit` module loaded.
 - Can access Blanca by loading `slurm/blanca` and submitting an interactive job.
- Summit
 - CU Boulder high performance computing resource.
 - Composed of forward-facing 483 compute nodes and 2 compile nodes.
- Blanca
 - Buy in condo nodes that compose a heterogeneous compute cluster.
 - Check with your group if you'd like access to this resource.

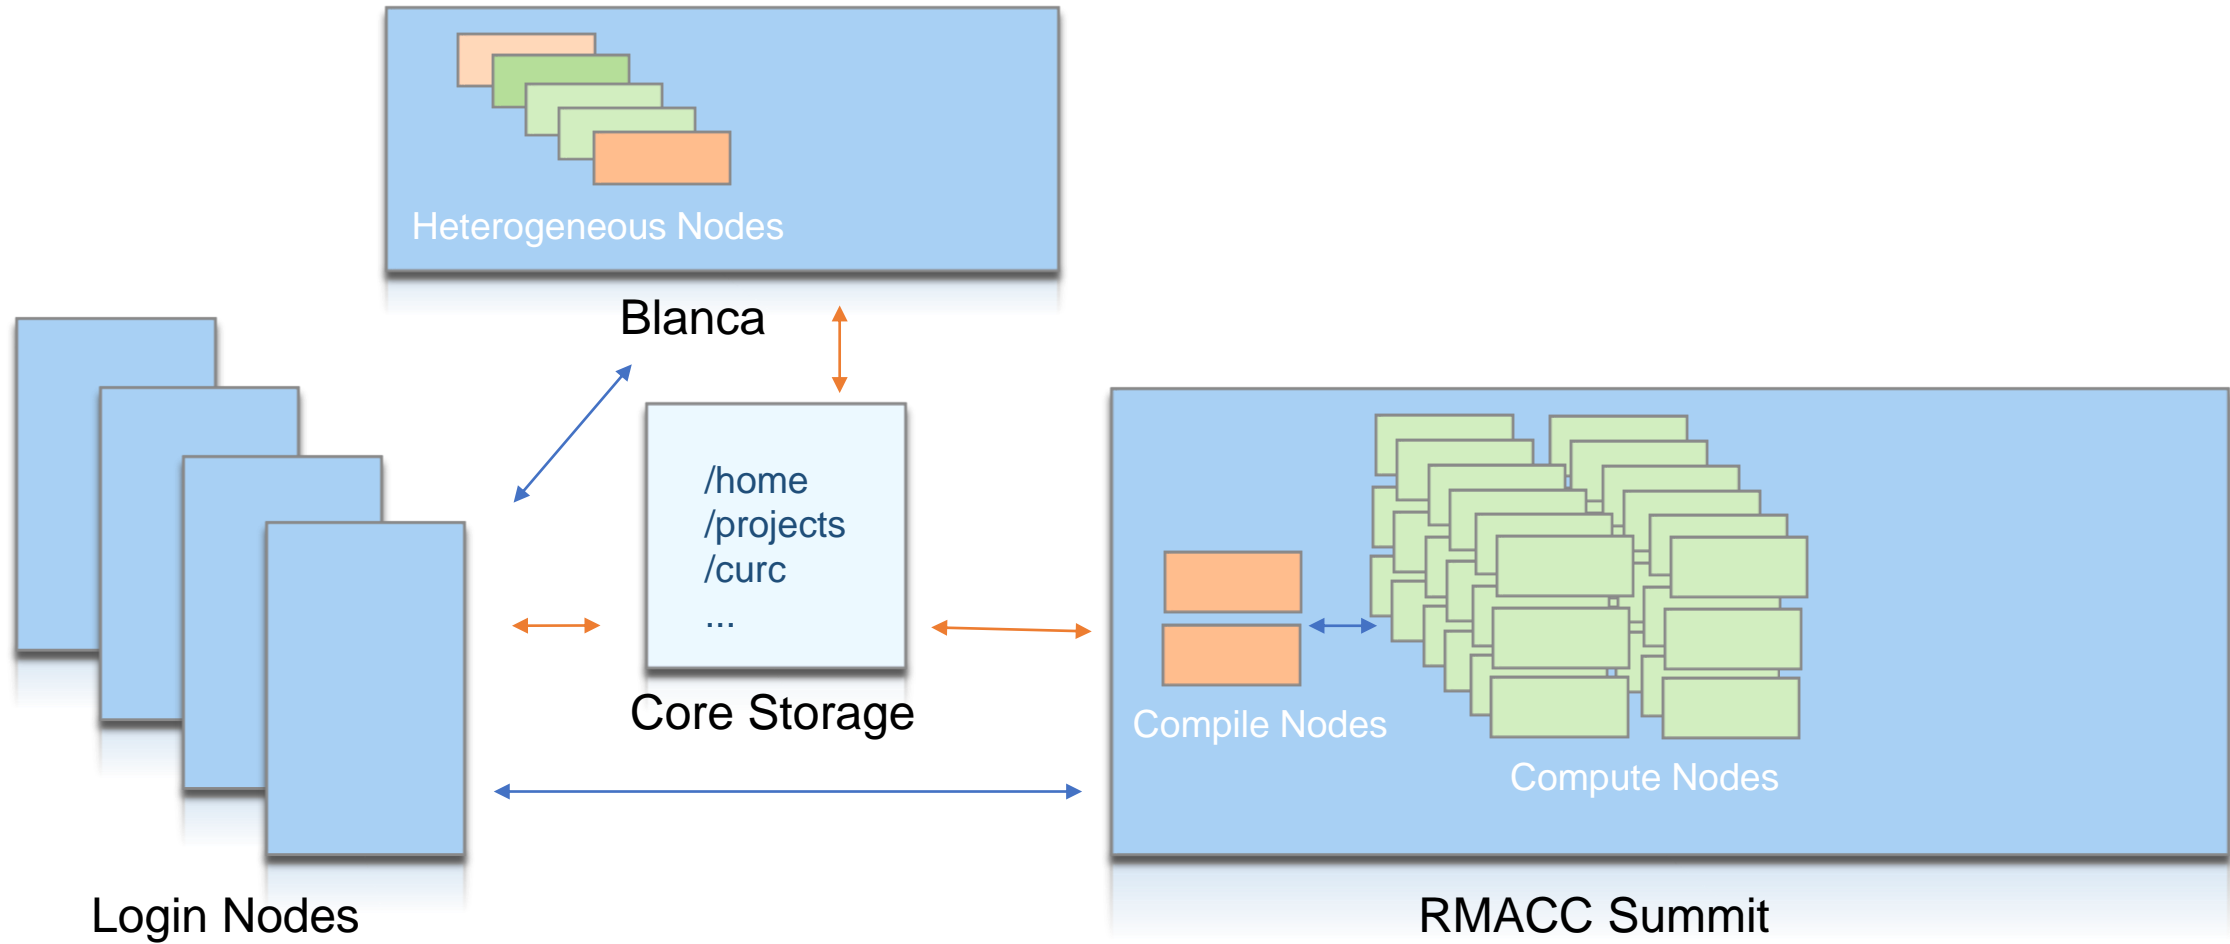
RC's network



RC's filesystems

- To reduce the amount of complexity for an end user to manage, RC uses a shared file server to manage user related storage.
- “Core Storage”
 - Contains `/home`, `/projects`, and `/rc_scratch`
 - Contains all shared software and the module stack.
- Every node or login VM is connected to this resource allow user to easily manage their files.
- Non-Parallel IO
- Slow reads and writes

RC's Storage



Parallel filesystems

- RC provides an additional system local parallel file system available on Summit.
- Scratch Space:
 - Spinning disk platters rated at 12 Gb/s
 - GPFS File System for parallel I/O w/ 32 Clients and 4 Servers
 - Locally mounted on Summit nodes.
 - Physically close to cores.
- Scratch is usually limited to 10 TB of storage but can be expanded upon with request.
 - Purges every 90 days from file creation
 - Technically shared among all users
 - Scratch totals in 1 PB of storage for everyone

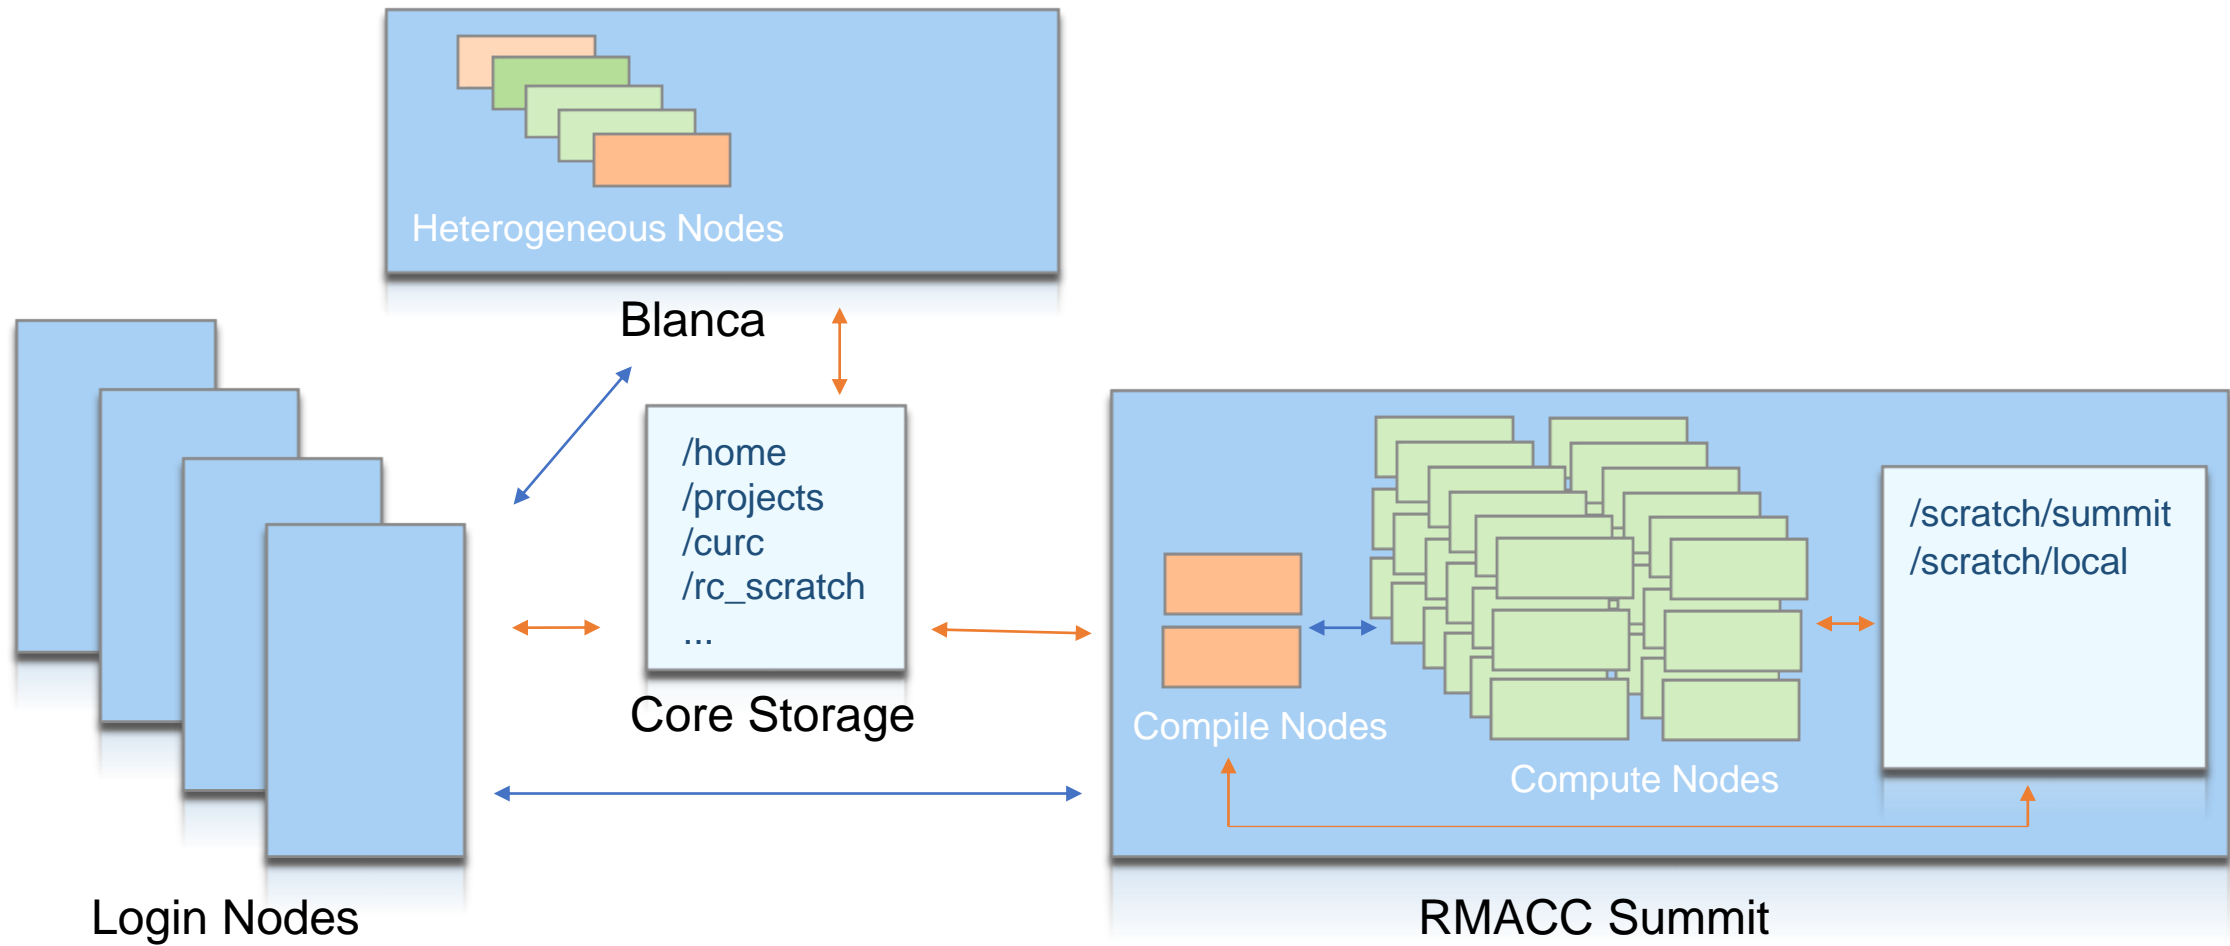
Local Node SSDs

- Summit Nodes also hold onto 100 GB of local node SSD storage.
- These SSDs are not shared among nodes so must move files over within job.
- No Cooperative Parallel I/O
- Located at `/scratch/local`

Some more filesystems...

- High Performance I/O is less available with Blanca
- Petalibrary Active
 - High Speed performance
 - Allows for parallel reads and writes
 - Currently utilizes BeeGFS as its parallel filesystem, will be switching to GPFS
 - More on this in a minute...
- RC_Scratch
 - Not the best storage on Summit...
 - Older scratch space past its service dates
 - Large but slow
 - *"I'm sometimes wonder if /projects or /home would be faster than rc_scratch..."*

Parallel Filesystems



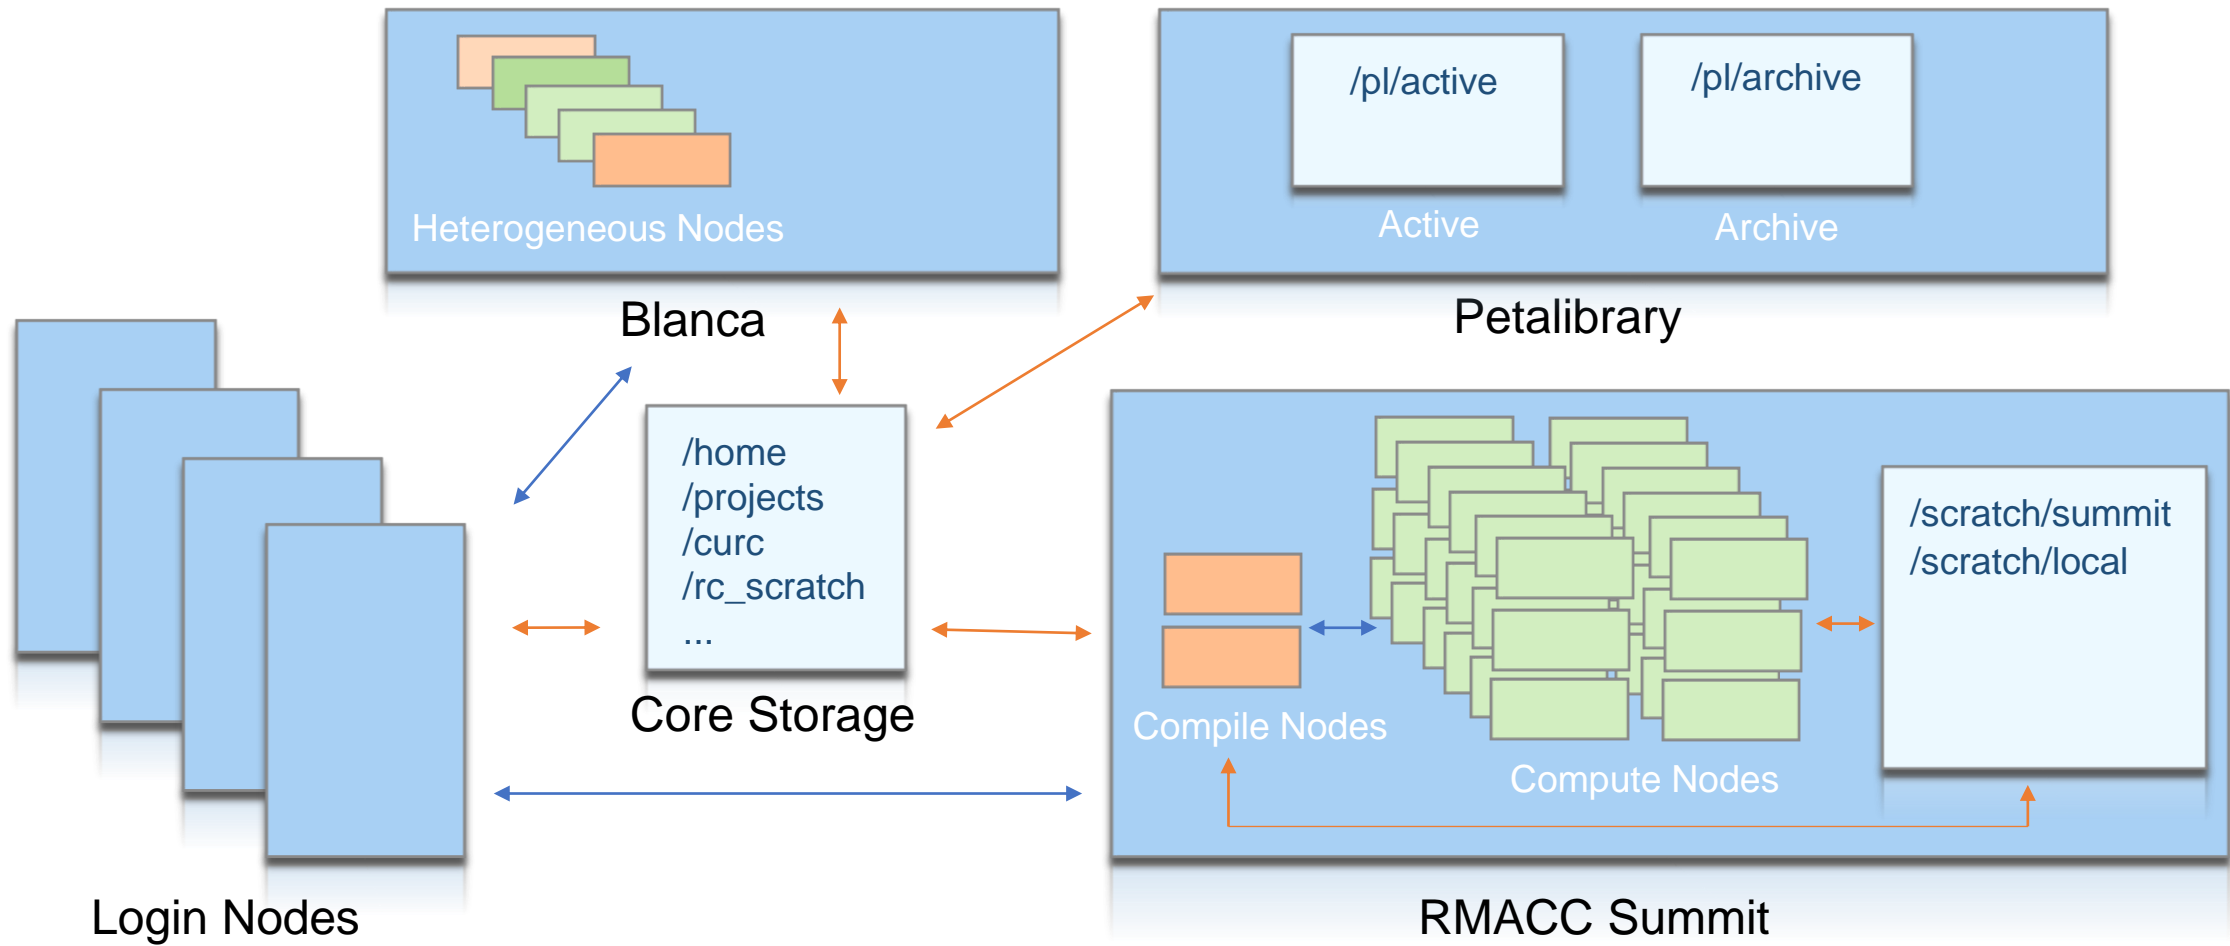
What is Petalibrary?

- Research Computing offers a subsidized but **paid**, long-term storage solution closely coupled with RC resources.
- Petalibrary
 - Large scale subsidized storage solution
 - Enterprise Grade
 - RC Staff supported with assistance on transfer strategies
 - Available in several flavors:
 - Active – Disk
 - Archival – Tape
 - Active Storage with Archive copy

Hardware Specifications

- Active Storage
 - Spinning disk platters for frequent reads and writes
 - BeeGFS filesystem
 - Parallel file I/O capable
 - RAID-6 file protection
 - Allocations located at: [/pl/active/](#)
- Archive Storage
 - Tape storage for infrequent reads and writes
 - iRods backed with StrongBox
 - Redundant copies of Data on separate tapes
 - Allocations located at: [/pl/archive/](#)

Petalibrary



Useful storage command:

- *curc-quota* – Research computing tool to monitor disk usage.
 - Provides detailed summary of your core storage
 - Provides detailed summary of scratch space on compile and compute nodes
 - Also lists current capacity of all Petalibrary allocations you have access to

```
[userXXXX@login12 ~]$ curc-quota
```

Data Transfers

- Data transfers are usually handled by one of 2 methods:
- Globus
 - By far the most stable and recommended way for data transfers
 - Fast transfers
 - Transfers continue if a user disconnects
 - Web GUI option or Globus Connect Personal
- SCP/SFTP
 - Secure Copy and Secure File Transfer Protocol
 - Straightforward method of transferring data
 - Generally recommend only to move small files less than a Gigabyte.



More on Data transfers

- Less common methods of transferring data...
- sshfs
 - Mounting the RC filesystem to your drive remotely!
 - Single sign in for multiple data transfers
 - Great when needing to repeatedly access files on RC Resources
- rsync/rclone
 - Another utility to transfer files
 - Particularly useful in repeated file transfers and synchronization of file sets
 - Snapshot like backups
 - <https://github.com/ResearchComputing/Documentation/blob/dev/docs/compute/rclone.md>

Some additional notes...

- Do not run your job against `/home` or `/projects`
 - Slows down your jobs.
 - Slows down everyone else's jobs.
 - We monitor for this and will kindly ask you to stop.
 - Use Scratch instead for your data!
- Always recover your data from scratch after your job completes!
 - No backup and we can't get your data back!
- Always have a safe backed up location for your data!

Questions?

Thank you!

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact information: rc-help@Colorado.edu
- Slides: https://github.com/ResearchComputing/Filesystems_And_Storage_Fall_2020